

# Introduction to Dynamic Programming

Soo Min Kwon

Department of Electrical and Computer Engineering  
Rutgers, The State University of New Jersey

September 10, 2021



# Outline

## 1 Deterministic Dynamic Programming

- Notations and Objective

- Example Problem

- Dynamic Programming Algorithm

- Approximation in Value Space

## 2 Stochastic Dynamic Programming

- DP Algorithm for Stochastic Problems



# Deterministic problem setting

A deterministic DP problem involves a discrete-time dynamic system that evolves over finite steps  $N$  and is of the form

$$x_{k+1} = f_k(x_k, u_k), \quad k = 0, 1, \dots, N - 1, \quad (1)$$

where

- $k$  is the time index
- $x_k$  is the state of the system (an element of some space)
- $u_k$  is the control or decision variable, to be selected at time  $k$  from some given set  $U_k(x_k)$  that depends on  $x_k$
- $f_k$  is a function of  $(x_k, u_k)$  that describes how the state is updated from time  $k$  to  $k + 1$
- $N$  is the horizon or number of times that control is applied




# Deterministic problem setting

Some more notation:

- *state space* is the set of all possible  $x_k$  at time  $k$
- *control space* is the set of all possible  $u_k$  at time  $k$
- $g_k(x_k, u_k)$  is the cost function incurred at time  $k$  that takes real number values
- $g_N(x_N)$  is a terminal cost incurred at the end of the process



# Example of deterministic DP



**Figure:** From state  $x_k$ , we can move to the next state under some nonrandom control  $u_k$  according to  $x_{k+1} = f_k(x_k, u_k)$



# Deterministic DP objective

- For a given initial state  $x_0$ , the total cost of control sequence  $\{u_0, \dots, u_{N-1}\}$  is


$$J(x_0; u_0, \dots, u_{N-1}) = g_N(x_N) + \sum_{k=0}^{N-1} g_k(x_k, u_k) \quad (2)$$

- We want to minimize this total cost over all control sequences to obtain the optimal value as a function of  $x_0$ :

$$J^*(x_0) = \min_{u_k \in U_k(x_k)} J(x_0; u_0, \dots, u_{N-1}) \quad (3)$$



# Example of a discrete control problem



- Example of a combinatorial finite-state, finite-horizon optimal control problem
- Goal is to produce a certain product from an initial state by performing a combination of operations A, B, C, and D
- Operation B can only be performed after A, and D can be performed only after C



# Principle of Optimality

The DP algorithm is based on the *principle of optimality*:

Let  $\{u_0^*, \dots, u_k^*, u_{k+1}^*, \dots, u_{N-1}^*\}$  be an optimal control sequence, which determines the corresponding optimal state sequence,  $\{x_1^*, \dots, x_{N-1}^*\}$ . Consider the subproblem in which we start at  $x_k^*$  at time  $k$  and wish to minimize the "cost-to-go" from time  $k$  to  $N$ ,

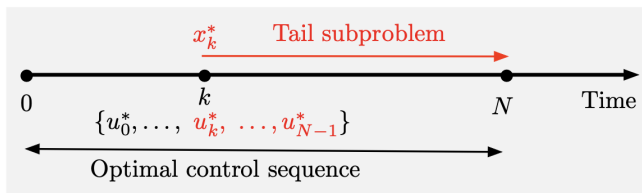
$$g_k(x_k^*, u_k) + \sum_{m=k+1}^{N-1} g_m(x_m, u_m) + g_N(x_N),$$

over  $\{u_k, \dots, u_{N-1}\}$  with  $u_m \in U_m(x_m)$ ,  $m = k, \dots, N-1$ . Then the truncated optimal control sequence  $\{u_k^*, \dots, u_{N-1}^*\}$  is optimal for the subproblem.





# Intuition behind the principle of optimality




**Figure:** The tail  $\{u_k^*, \dots, u_{N-1}^*\}$  of an optimal control sequence  $\{u_0^*, \dots, u_k^*, u_{k+1}^*, \dots, u_{N-1}^*\}$  is optimal for the tail subproblem that starts at  $x_k^*$

- **Simple analogy:** If the fastest route from Los Angeles to Boston passes through Chicago, the principle of optimality states that the Chicago to Boston portion of the route is also the fastest for a trip that starts in Chicago and ends in Boston.



# Solving the previous example



Solution:

- Solve the tail subproblem of length 2, then 3, then 4
- Note the shortest path from the initial state to the terminal state



# DP algorithm for deterministic finite horizon problems

With the previous example, we can now construct the DP algorithm:

Start with

$$J_N^*(x_N) = g_N(x_N), \quad \text{for all } x_N, \quad (4)$$

and for  $k = 0, \dots, N - 1$ , let

$$J_k^*(x_k) = \min_{u_k \in U_k(x_k)} [g_k(x_k, u_k) + J_{k+1}^*(f_k(x_k, u_k))], \quad (5)$$

for all  $x_k$ .

- The algorithm is solving every tail subproblem by constructing functions  $J_N^*(x_N), \dots, J_0^*(x_0)$
- The last step gives us the optimal cost  $J^*(x_0)$



# Constructing the optimal control sequence

With the previous algorithm, we can obtain the functions  $J_0^*, \dots, J_N^*$  and solve for the optimal control sequence:

**Construction of Optimal Control Sequence  $\{u_0^*, \dots, u_{N-1}^*\}$**

Set

$$u_0^* \in \arg \min_{u_0 \in U_0(x_0)} \left[ g_0(x_0, u_0) + J_1^*(f_0(x_0, u_0)) \right],$$

and

$$x_1^* = f_0(x_0, u_0^*).$$

Sequentially, going forward, for  $k = 1, 2, \dots, N-1$ , set

$$u_k^* \in \arg \min_{u_k \in U_k(x_k^*)} \left[ g_k(x_k^*, u_k) + J_{k+1}^*(f_k(x_k^*, u_k)) \right], \quad (1.7)$$

and

$$x_{k+1}^* = f_k(x_k^*, u_k^*). \quad (1.8)$$

**Figure:** Construction of Optimal Control Sequence



# Approximation in value space

- Of course, in practice, computing  $J_k^*(x_k)$  can be expensive
- Instead, find an approximation  $\tilde{J}_k$  for a suboptimal solution

## Approximation in Value Space - Use of $\tilde{J}_k$ in Place of $J_k^*$

Start with

$$\tilde{u}_0 \in \arg \min_{u_0 \in U_0(x_0)} \left[ g_0(x_0, u_0) + \tilde{J}_1(f_0(x_0, u_0)) \right],$$

and set

$$\tilde{x}_1 = f_0(x_0, \tilde{u}_0).$$

Sequentially, going forward, for  $k = 1, 2, \dots, N - 1$ , set

$$\tilde{u}_k \in \arg \min_{u_k \in U_k(\tilde{x}_k)} \left[ g_k(\tilde{x}_k, u_k) + \tilde{J}_{k+1}(f_k(\tilde{x}_k, u_k)) \right], \quad (1.9)$$

and

$$\tilde{x}_{k+1} = f_k(\tilde{x}_k, \tilde{u}_k). \quad (1.10)$$

Figure: Approximation in Value Space



# Q-Factors

The minimization expression of equation (1.9) is the **Q-factor** of  $(x_k, u_k)$ :

$$\tilde{Q}_k(x_k, u_k) = g_k(x_k, u_k) + \tilde{J}_{k+1}(f_k(x_k, u_k)) \quad (6)$$



# Stochastic dynamic programming

- Stochastic DP problems have the form

$$x_{k+1} = f_k(x_k, u_k, w_k), \quad k = 0, 1, \dots, N-1, \quad (7)$$

where  $w_k$  is a "disturbance" (e.g. physical noise, market uncertainties, demand for inventory) characterized by a probability distribution  $P_k(\cdot|x_k, u_k)$  that may depend on  $x_k, u_k$ , but not on prior disturbances

- An **important** difference is that we optimize not over the control sequences, but rather over *policies*

$$\pi = \{\mu_0, \dots, \mu_{N-1}\},$$

where  $u_k = \mu_k(x_k)$ .



# Stochastic dynamic programming

- In the presence of uncertainty, optimizing over the policies can give us improved costs, as they allow choices that use knowledge of  $x_k$
- Given an initial state  $x_0$  and policies  $\pi = \{\mu_0, \dots, \mu_{N-1}\}$ , the expected cost of  $\pi$  starting at  $x_0$  is

$$J_\pi(x_0) = E \left\{ g_N(x_N) + \sum_{k=0}^{N-1} g_k(x_k, \mu_k(x_k), w_k) \right\}$$

- An optimal policy  $\pi^*$  is one that minimizes the cost

$$J_{\pi^*}(x_0) = \min_{\pi \in \Pi} J_\pi(x_0)$$





# DP algorithm for stochastic finite horizon problems

Similar to that of the deterministic algorithm except with random variables:

## DP Algorithm for Stochastic Finite Horizon Problems

Start with

$$J_N^*(x_N) = g_N(x_N), \quad (1.12)$$

and for  $k = 0, \dots, N-1$ , let

$$J_k^*(x_k) = \min_{u_k \in U_k(x_k)} E \left\{ g_k(x_k, u_k, w_k) + J_{k+1}^*(f_k(x_k, u_k, w_k)) \right\}. \quad (1.13)$$

If  $u_k^* = \mu_k^*(x_k)$  minimizes the right side of this equation for each  $x_k$  and  $k$ , the policy  $\pi^* = \{\mu_0^*, \dots, \mu_{N-1}^*\}$  is optimal.

